

# Classification Progress

Machine Learning Project

November 2025

Team Members:  
Naol Seyum  
Emery Hagye

# Github Creation with Added Features



- Created a simple Github Repository for keeping progress
- Used Git Actions to create a Tiny CL
- Implemented a pre-commit to help further with formatting

The screenshot shows a GitHub repository page for a private project named "Machine-Learning-Classification-Project". The repository has 1 branch and 0 tags. The main branch contains 15 commits from the user "ehagye". The commits are:

- Merging data-loader to main branch (80610ef · 2 weeks ago)
- .github/workflows (chore(tiny ci): attempting to get the ci working again without...) (3 weeks ago)
- src (these are apart of the last commit as well) (3 weeks ago)
- .gitignore (feat(data\_loader): Created a data loader for the txt files give...) (3 weeks ago)
- .pre-commit-config.yaml (added automatic code quality with black and ruff via pre-co...) (3 weeks ago)
- README.md (Update README.md) (3 weeks ago)
- pyproject.toml (fixed formating of the .toml file for the ruff tool as it requires...) (3 weeks ago)
- requirements.txt (Added a requirements.txt file and have VS Code enviroment ...) (3 weeks ago)

The repository has 0 forks, 0 stars, and 0 watching. It also has 0 releases and 0 packages published. The Languages section shows Python at 100.0%.

**Machine Learning Classification Project**

# Data Loader Python File

- Created a Data Loader that Automatically finds correct path
- Can easily be reused anywhere in the project
- A little preprocessing with replacing the missing entries with nan

```
src > data_loader.py > load_dataset
1  from pathlib import Path
2
3  import numpy as np
4
5  missing_entry = 1e99
6
7  def _project_root() -> Path:
8      return Path(__file__).resolve().parents[1]
9
10 def _resolve(pathlike) -> Path:
11     p = Path(pathlike)
12     if not p.is_absolute():
13         p = _project_root() / p
14     return p
15
16 def load_dataset(train_path, label_path=None, test_path=None):
17     """
18         Loads training/test datasets and changes the missing values (1e99) to be Nan.
19
20     Parameters
21     -----
22     train_path: string
23         this is a path to the training data file.
24     label_path: string or none
25         this is a path to the training labels file.
26     test_path: string or none
27         this is a path to the test data file.
28
29
30
31     Returns
32     -----
33     dict type
34         A dictionary with the keys X_train, Y_train, and X_test.
35         (only when they are given as parameters)
36
37     # loading the data from the given files and replacing the missing entries with NaN
38     def load(path):
39         path = _resolve(path)
40
41         arr = np.loadtxt(path)
42         arr = np.where(arr == missing_entry, np.nan, arr)
43         return arr
44
45     data = {"X_train": load(train_path)}
46     if label_path:
47         data["Y_train"] = load(label_path).astype(int).ravel()
48     if test_path:
49         data["X_test"] = load(test_path)
50
51     return data["X_train"], data.get("Y_train"), data.get("X_test")
```

# Began EDA of Dataset 1



```
Total Nans: 9,936
```

```
Top 10 features by most Nans:
```

```
f1    12
```

```
f1    10
```

```
f1     9
```

```
f1     8
```

```
dtype: int64
```

```
Top 10 rows with most missing values:
```

```
58     85
```

```
51     85
```

```
89     83
```

```
109    82
```

```
140    81
```

```
47     80
```

```
73     79
```


```
133    79
```

```
117    78
```

```
71     78
```

```
dtype: int64
```

Class counts (training)



# Methods to use in the future



- For Datasets 1 and 2 I will use aggressive Dimensionality reduction for the large number of features
  - Variance thresholding
  - PCA
- Complex Model Selection
  - Linear SVM
  - Random Forest
  - Neural Network



# Spam Email Detection

Machine Learning Project

November 2025


Team Members:  
Naol Seyum  
Emery Hagye

# Why Spam Detection?

- Inboxes overflow with spam and scams.



- ML can detect spam automatically.



# Project Workflow



Emails



Preprocessing



Feature Extraction




Feature Extraction



Predictions

- Combined two training datasets.
- Cleaned and vectorized text.
- Trained an SVM model.
- Tuned with cross-validation.
- Predicted spam or ham on test data.



```

9 import pandas as pd
10
11 # step-1 data loading part
12 # this part just reads the files and makes some columns
13 train1 = pd.read_csv("spam_train1.csv")      # columns v1 (label), v2 (text),
14 train2 = pd.read_csv("spam_train2.csv")      # columns label, text, label_num, etc.
15 test   = pd.read_csv("spam_test.csv")        # column message
16
17 # step-2 cleaning up the data
18 # just renaming and merging both train sets so everything looks the same.
19 train1_simple = train1[['v1', 'v2']].rename(columns={'v1': 'label', 'v2': 'text'})
20 train2_simple = train2[['label', 'text']]
21 data = pd.concat([train1_simple, train2_simple], ignore_index=True)
22
23 # step-3 turning labels into numbers.
24 # ham = 0, spam = 1 so the model can understand it.
25 data['label_num'] = data['label'].map({'ham': 0, 'spam': 1})
26

```

```

27 # step-4 feature setup
28 # hashing vectorizer turns words into numbers, linear SVM does the classification.
29 from sklearn.feature_extraction.text import HashingVectorizer
30 from sklearn.svm import LinearSVC
31 from sklearn.pipeline import make_pipeline
32
33 pipe = make_pipeline(
34     HashingVectorizer(
35         n_features=2**18,
36         alternate_sign=False,
37         ngram_range=(1, 2) # unigrams + bigrams
38     ),
39     LinearSVC()
40 )

```

# Data Preprocessing

Merging two training datasets.

- Merged two training datasets into one.

Encoded labels

- ham = 0
- spam = 1

# Hashing Vectorizer

Renamed columns

- v1 → label
- v2 → text

- Converted text into numeric features

# Training the Model & Checking Performance


- Split data = 80% for training and 20% for validation.
- Trained Support Vector Machine (SVM).
- Checked Accuracy, Precision, Recall, and F1-score.
- Used GridSearchCV for tuning ( $C = [0.1, 1, 10]$ ).

```
42 # step-5 training and validation split.  
43 | # splitting the data so I can test the model's performance before final training.  
44 from sklearn.model_selection import train_test_split  
45  
46 X_train, X_val, y_train, y_val = train_test_split(  
47 | data['text'], data['label_num'],  
48 | test_size=0.2,  
49 | stratify=data['label_num'],  
50 | random_state=42  
51 )  
52 # step-6 training the model and checking accuracy.  
53 | # fitting the model and then printing accuracy, precision, recall, and f1.  
54 from sklearn.metrics import accuracy_score, precision_recall_fscore_support, roc_auc_score  
55  
56 pipe.fit(X_train, y_train)  
57 y_pred = pipe.predict(X_val)  
58 acc = accuracy_score(y_val, y_pred)  
59 prec, rec, f1, _ = precision_recall_fscore_support(y_val, y_pred, average='binary')  
60  
61 print("Baseline LinearSVC + HashingVectorizer")  
62 print(f"Accuracy: {acc:.4f}")  
63 print(f"Precision: {prec:.4f}")  
64 print(f"Recall: {rec:.4f}")  
65 print(f"F1-score: {f1:.4f}")  
66 # step-7 tuning time.  
67 | # trying different C values to see which one gives the best spam detection score  
68 from sklearn.model_selection import GridSearchCV  
69  
70 param_grid = [  
71 | 'linearsvc_C': [0.1, 1, 10]  
72 ]  
73 grid = GridSearchCV(  
74 | pipe,  
75 | param_grid=param_grid,  
76 | cv=5,  
77 | scoring='f1' # bias towards good spam detection.  
78 )  
79  
80 grid.fit(data['text'], data['label_num'])  
81 best_model = grid.best_estimator_  
82  
83 print("Best C from GridSearchCV:", grid.best_params_)
```

**True Positive (TP)** → spam predicted as spam.

**False Positive (FP)** → ham predicted as spam.

**False Negative (FN)** → spam predicted as ham.



```
85 # step-8 final training and test prediction
86 # model done training, now predicting on the test file and saving it to csv
87 best_model.fit(data['text'], data['label_num'])
88
89 test_text = test['message']
90 test_preds = best_model.predict(test_text)
91
92 # step-9 mapping numbers back to labels
93 # just converting 0 and 1 back to ham and spam so it's readable
94 label_map = {0: 'ham', 1: 'spam'}
95 pred_labels = [label_map[p] for p in test_preds]
96
97 # step-10 saving the result
98 # final output with id and predicted label
99 output = pd.DataFrame({
100     'id': range(len(test)),
101     'label': pred_labels
102 })
103
104 output.to_csv("spam_test_predictions.csv", index=False)
105 print("The prediction result is saved to spam_test_predictions.csv")
106
107 # step-11 quick summary of prediction counts
108 # just counting how many are spam vs ham in the final predictions
109 ham_count = (output['label'] == 'ham').sum()
110 spam_count = (output['label'] == 'spam').sum()
111
112 print(f"\nSummary of predictions:")
113 print(f"Ham emails : {ham_count}")
114 print(f"Spam emails: {spam_count}")
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS zsh + ⌂ ⌄ ⌓ ⌔ ⌕ ⌖

```
naolseyum@Mac Spam Email Detection 2 % python3 ham_spam_email_detection.py
Baseline LinearSVC + HashingVectorizer
Accuracy: 0.955
Precision: 0.9620
Recall: 0.8172
F1 score: 0.8900
Best C from GridSearchCV: {'linear_svc__C': 10}
The prediction result is saved to spam_test_predictions.csv

Summary of predictions:
Ham emails : 5248
Spam emails: 1199
```

Model accuracy	Precision	Best parameter	Ham: X		Spam: Y
----------------	-----------	----------------	--------	--	---------

• 95.3%	• 96.2%	• C=10	• 5248	• 1199
---------	---------	--------	--------	--------

- Predictions saved to: `spam_test_predictions.csv`

# Model Performance & Output Summary